

# Brandon Yifan Yang

[brandonyifanyang.com](http://brandonyifanyang.com)   [yang52@engineering.upenn.edu](mailto:yang52@engineering.upenn.edu)   [github.com/branyang02](https://github.com/branyang02)   [Google Scholar](#)

## EDUCATION

---

- **University of Pennsylvania** Philadelphia, PA  
*M.S.E. in Robotics; GPA: 3.7/4.0* May 2027
- **University of Virginia** Charlottesville, VA  
*B.S. in Computer Science; Highest Distinction; GPA: 3.9/4.0* May 2025
- **Selected Coursework:** Learning in Robotics, Reinforcement Learning, Probabilistic Machine Learning, GPU Programming & Architecture, Physical Intelligence, Natural Language Processing

## EXPERIENCE

---

**Research Lead** San Francisco, CA  
*AfterQuery (YC W25)* May 2026 - Present

- Led development of custom UMI hardware and a distributed robot data collection system, integrating Rust-based device-side capture and LiDAR tracking with a server-side SLAM pipeline for end-effector tracking and automated storage with Hugging Face dataset management.
- Led a 6-person research team building an AI-agent search benchmark, including task collection and validation, rubric grading, agent-run orchestration, trace logging, and reproducible model comparisons.

**Embodied Intelligence Research Intern** Beijing, China  
*Spirit AI* May 2025 - Aug 2025

- Trained and evaluated VLA policy variants in PyTorch across simulation and real dual-arm robot experiments, comparing contrastive objectives, observation-noise curricula, reduced guidance inputs, and mixture-of-experts architectures.
- Designed and implemented a DAgger data-collection pipeline for VLA models, including operator-takeover logic and corrective-demonstration workflows used for model training and evaluation.
- Built a reproducible Docker/Slurm pipeline for VLA simulation, training, and evaluation workflows; open-sourced code and datasets on [GitHub](#).

## RESEARCH EXPERIENCE

---

**Perception, Action, & Learning Lab, University of Pennsylvania** Philadelphia, PA  
*Advisor: Dinesh Jayaraman* Aug 2025 - Present

- Implemented demonstration retrieval for  $\pi_{0.5}$  VLA policies, integrating robot data retrieval into training and test-time inference for tabletop manipulation experiments. (*CoRL 2026*)
- Engineered data transformation, training, and evaluation pipelines for retrieval-based VLA variants using JAX, PyTorch, Hugging Face, Docker, Slurm, and Ray.
- Designed activation-collection and activation-steering infrastructure for VLA analysis across MetaWorld, LIBERO, RoboCasa, and real Franka hardware experiments. (*ICML 2026 Workshop; NeurIPS 2026*)

**Learning and Interactive Robotics, University of Virginia** Charlottesville, VA  
*Advisor: Yen-Ling Kuo* Aug 2024 - May 2025

- **SkillVLA:** Developed a skill-conditioned VLA model for language-conditioned manipulation, using subgoal instructions and a learned skill library to structure action generation.
- Presented SkillVLA as an oral talk at the 2024 UVA LLM Workshop ([slides](#)), earning the Audience Choice Award (top 3 of 28 presentations).
- **Contrastive Learning for Robot Manipulation:** Trained contrastive objectives over action sequences to learn behavior-grounded visual embeddings for visuomotor policies under varied camera poses and object appearances. ([CoRL 2025](#))
- Modified simulation environments and trained PyTorch visuomotor policies to evaluate learned behavior-grounded embeddings.

- **Grounded Location for Object Manipulation (GLOMA)** [code]: Led a team of 3 in developing a language-grounded image-editing pipeline for object relocation goals in robotic manipulation settings.
- Integrated language grounding with bounding-box-guided visual editing for object relocation.
- Collected and annotated a dataset for fine-tuning pre-trained language and vision models.
- **Centralized Multi-Agent RL for Collaborative Tasks** [code]: Implemented offline and online centralized MARL training for long-horizon robotic bolt-screwing tasks in Isaac Gym.
- Designed custom multi-agent reward functions for task completion and collaboration, improving simulated task success by 20% over the project baseline.

## PUBLICATIONS

---

Jiani Huang\*, **Brandon Y. Yang\***, Amish Sethi, Yuchen Zheng, Christopher Watson, Jianing Qian, Mayur Naik, and Dinesh Jayaraman. “Retrieval-Augmented Vision-Language-Action Model.” Under review at *CoRL 2026*.

Miranda Muqing Miao\*, Subin Kim\*, **Brandon Y. Yang\***, Dinesh Jayaraman, and Lyle Ungar. “Contrastive Conceptor Activation Steering (COAST): Unlocking Vision-Language-Action Models through Hidden States.” Accepted poster, *ICML 2026 Mechanistic Interpretability Workshop*; submitted to *NeurIPS 2026*.  
[Workshop][arXiv]

Amish Sethi\*, **Brandon Y. Yang\***, Michael Evans, Spencer Mateega, and Sam Jung. “Position: Trust No Multi-Turn Benchmark Without Measuring Its User Simulator.” Submitted to *NeurIPS 2026 Position Paper Track*.

Sung-Wook Lee, Xuhui Kang, **Brandon Y. Yang**, and Yen-Ling Kuo. “CLASS: Contrastive Learning via Action Sequence Supervision for Robot Manipulation.” *Conference on Robot Learning*, pp. 4743–4766. PMLR, 2025.  
[Website][arXiv][PDF]

\*Equal contribution among starred authors.

## SOFTWARE PROJECTS

---

**openpi-eval** [GitHub, Dataset]: Built OpenPI evaluation infrastructure for a NeurIPS submission, supporting reproducible pi05 rollouts across MetaWorld, LIBERO, LIBERO-Pro, RoboCasa, and RoboLab, plus pi0-fast support where released and server-side inference batching.

**openpi-cuda** [GitHub]: Developed custom CUDA kernels with C++/Python bindings for  $\pi_{0.5}$  VLA inference, reducing latency by 18.5ms versus the PyTorch baseline.

**notie-markdown** [Website]: Built a React/TypeScript markdown renderer with equation previewing, graphing, and code execution; shipped with a Python/Flask backend on Vercel and Heroku.

**SmartOH** [GitHub]: Built an office-hours queue management system with real-time queue updates, notifications, and analytics using Next.js, TypeScript, Python, and FastAPI; deployed with Vercel, AWS EC2, and CI/CD.

## SKILLS

---

**Programming:** Python, C++, CUDA C++, TypeScript/JavaScript, Java

**ML & Robotics:** PyTorch, JAX, Hugging Face, CUDA, Ray, Slurm, robot learning, VLA models

**Web & Infrastructure:** React, FastAPI, Flask, Docker, AWS EC2, Vercel, uv, ruff